

Intro to R

Functions, Objects and Visualization

Prof. Bisbee

Seoul National University

Slides Updated: 2024-06-30

Agenda

1. Recap of last lecture

- Using packages: `install.packages()` & `require()`
- Loading and manipulating data: `read_rds()` and `%>%`

2. `tidyverse` functions

- `filter` and `select`
- `summarize` and `mutate`
- `group_by`

Loading Packages & Data

- Create an `.Rmd` file and save to your `code` folder
 - Accept defaults, Save As... (with a good name), then `knit`
- Load the `tidyverse` package

```
require(tidyverse)
```

- Load the data from the course [github page](#) directly using `read_rds()`
 - We **create** an "object" to store the data using a left-arrow: `<-`

```
df<-  
read_rds("https://github.com/jbisbee1/ISP_Data_Science_2024/raw/main/da
```

Tabular Data

- Data comes in many different formats
- **Structured data:** standardized, well-defined structure, easily accessed
 - I.e., tables, databases
 - In my YouTube example, the survey we gave was **structured**
- **Unstructured data:** messy, organic, disorganized, hard to use
 - I.e., web pages, images, videos
 - In my YouTube example, the scraped HTML code of a list of recommendations was **unstructured**
- In this class, we will always be working with **structured** data...specifically "tabular data frames"
- This still requires work to prepare!

Tabular Data Frame

- AKA a "tibble"
- These are "square" (although actually rectangular)
- Rows: **units of observation** (i.e., the entities we are studying)
 - People (each row is a survey respondent, athlete, etc.)
 - Places (each row is a state, county, country, etc.)
 - Things (each row is a tweet, firm, product, etc.)
- Columns: **variables of interest** (i.e., attributes we are studying)
 - Beliefs / behaviors / etc. (i.e., where rows are people)
 - Rainfall / crimes / etc. (i.e., where rows are places)
 - Likes / profits / etc. (i.e., where rows are things)

Looking at Data

- We now have the contents of `sc_debt.Rds` stored in the object `df`
- We can look at this object directly

```
df
```

```
## # A tibble: 2,546 × 16
##   unitid instnm stabbr grad_debt_mdn control region preddeg
##   <int> <chr>   <chr>         <int> <chr>   <chr>   <chr>
## 1 100654 Alaba... AL             33375 Public  South... Bachel...
## 2 100663 Unive... AL             22500 Public  South... Bachel...
## 3 100690 Amrid... AL             27334 Private South... Associ...
## 4 100706 Unive... AL             21607 Public  South... Bachel...
## 5 100724 Alaba... AL             32000 Public  South... Bachel...
## 6 100751 The U... AL             23250 Public  South... Bachel...
## 7 100760 Centr... AL             12500 Public  South... Associ...
## 8 100812 Athen... AL             19500 Public  South... Bachel...
## 9 100830 Aubur... AL             24826 Public  South... Bachel...
## 10 100858 Aubur... AL             21281 Public  South... Bachel...
## # i 2,536 more rows
## # i 9 more variables: openadmp <int>, adm_rate <dbl>,
## #   ccbasic <int>, sat avg <int>, md earn wne p6 <int>,
```

Looking at Data

- What is our **unit of observation**?
 - Academic institutions: each row is a single school
- What are our **variables of interest**?
 - Let's look!

```
colnames(df) # Prints the variable names
```

```
## [1] "unitid"      "instnm"      "stabbr"  
## [4] "grad_debt_mdn" "control"     "region"  
## [7] "preddeg"     "openadmp"    "adm_rate"  
## [10] "ccbasic"     "sat_avg"     "md_earn_wne_p6"  
## [13] "ugds"        "costt4_a"    "selective"  
## [16] "research_u"
```

Good Data has Codebooks!

Name	Definition
unitid	Unit ID
instnm	Institution Name
stabbr	State Abbreviation
grad_debt_mdn	Median Debt of Graduates
control	Control Public or Private
region	Census Region
preddeg	Predominant Degree Offered: Associates or Bachelors
openadmp	Open Admissions Policy: 1=Yes, 2=No, 3=No 1st time students
adm_rate	Admissions Rate: proportion of applications accepted
ccbasic	Type of institution*
sat_avg	Average SAT scores
md_earn_wne_p6	Average Earnings of Recent Graduates
ugds	Number of undergraduates
costt4_a	Average cost of attendance (tuition-grants)
selective	Institution admits fewer than 10% of applications, 1=Yes, 0=No
research_u	Institution is a research university, 1=Yes, 0=No

Looking at data

- Looking at data is **crucial**

```
# First 6 rows  
df %>% head()
```

```
## # A tibble: 6 × 16  
##   unitid instnm  stabbr grad_debt_mdn control region preddeg  
##   <int> <chr>   <chr>      <int> <chr>   <chr> <chr>  
## 1 100654 Alabam... AL          33375 Public  South... Bachel...  
## 2 100663 Univer... AL          22500 Public  South... Bachel...  
## 3 100690 Amridg... AL          27334 Private South... Associ...  
## 4 100706 Univer... AL          21607 Public  South... Bachel...  
## 5 100724 Alabam... AL          32000 Public  South... Bachel...  
## 6 100751 The Un... AL          23250 Public  South... Bachel...  
## # i 9 more variables: openadmp <int>, adm_rate <dbl>,  
## #   ccbasic <int>, sat_avg <int>, md_earn_wne_p6 <int>,  
## #   ugds <int>, costt4_a <int>, selective <dbl>,  
## #   research_u <dbl>
```

- (Same as `head(df)`)

Looking at data

- Looking at data is **crucial**

```
# Last 6 rows  
df %>% tail()
```

```
## # A tibble: 6 × 16  
##   unitid instnm  stabbr grad_debt_mdn control region preddeg  
##   <int> <chr>   <chr>         <int> <chr>   <chr> <chr>  
## 1 493716 Yeshiv... NJ              NA Private North... Associ...  
## 2 493725 Univer... AR              NA Public  South... Bachel...  
## 3 493822 Colleg... RI              NA Private New E... Bachel...  
## 4 494630 Christ... TX              NA Private South... Bachel...  
## 5 494685 Urshan... MO              NA Private Plains Bachel...  
## 6 494737 Yeshiv... NY              NA Private North... Bachel...  
## # i 9 more variables: openadmp <int>, adm_rate <dbl>,  
## #   ccbasic <int>, sat_avg <int>, md_earn_wne_p6 <int>,  
## #   ugds <int>, costt4_a <int>, selective <dbl>,  
## #   research_u <dbl>
```

- (Same as `tail(df)`)

Manipulating the Data

- Last lecture, we wanted to know...

1. Where is **Vanderbilt University**?

```
df %>%  
  filter(instnm == "Vanderbilt University") # Only select rows with  
  Vandy
```

```
## # A tibble: 1 × 16  
##   unitid instnm  stabbr grad_debt_mdn control region preddeg  
##   <int> <chr>   <chr>         <int> <chr>   <chr> <chr>  
## 1 221999 Vander... TN             14962 Private South... Bachel...  
## # i 9 more variables: openadmp <int>, adm_rate <dbl>,  
## #   ccbasic <int>, sat_avg <int>, md_earn_wne_p6 <int>,  
## #   ugds <int>, costt4_a <int>, selective <dbl>,  
## #   research_u <dbl>
```

Manipulating the Data

- What if we don't know precisely how Vandy is spelled in these data?
- `str_detect()` and `grepl()` to the rescue!

```
df %>%  
  filter(str_detect(instnm, 'Vand'))
```

```
## # A tibble: 2 × 16  
##   unitid instnm  stabbr grad_debt_mdn control region preddeg  
##   <int> <chr>   <chr>          <int> <chr>   <chr> <chr>  
## 1 149639 Vander... IL          27000 Private Great... Bachel...  
## 2 221999 Vander... TN          14962 Private South... Bachel...  
## # i 9 more variables: openadmp <int>, adm_rate <dbl>,  
## #   ccbasic <int>, sat_avg <int>, md_earn_wne_p6 <int>,  
## #   ugds <int>, costt4_a <int>, selective <dbl>,  
## #   research_u <dbl>
```

Manipulating the Data

- What if we don't know precisely how Vandy is spelled in these data?
- `str_detect()` and `grepl()` to the rescue!

```
df %>%  
  filter(grepl('Vand', instnm))
```

```
## # A tibble: 2 × 16  
##   unitid instnm  stabbr grad_debt_mdn control region preddeg  
##   <int> <chr>   <chr>          <int> <chr>   <chr> <chr>  
## 1 149639 Vander... IL          27000 Private Great... Bachel...  
## 2 221999 Vander... TN          14962 Private South... Bachel...  
## # i 9 more variables: openadmp <int>, adm_rate <dbl>,  
## #   ccbasic <int>, sat_avg <int>, md_earn_wne_p6 <int>,  
## #   ugds <int>, costt4_a <int>, selective <dbl>,  
## #   research_u <dbl>
```

Manipulating the Data

- We can go deeper with this logic
 - "or" denoted with `|`
 - "and" denoted with `&`

```
df %>%  
  filter(str_detect(instnm, "Vand") | str_detect(instnm, "Tenn"))
```

```
## # A tibble: 12 × 16  
##   unitid instnm stabbr grad_debt_mdn control region preddeg  
##   <int> <chr>   <chr>      <int> <chr>   <chr>   <chr>  
## 1 149639 Vande... IL          27000 Private Great... Bachel...  
## 2 220075 East ... TN          20500 Public  South... Bachel...  
## 3 220978 Middl... TN          21500 Public  South... Bachel...  
## 4 221485 South... TN           NA Public  South... Associ...  
## 5 221731 Tenne... TN          21500 Private South... Bachel...  
## 6 221740 The U... TN          20635 Public  South... Bachel...  
## 7 221759 The U... TN          20500 Public  South... Bachel...  
## 8 221768 The U... TN          22500 Public  South... Bachel...  
## 9 221838 Tenne... TN          27000 Public  South... Bachel...  
## 10 221847 Tenne... TN          17000 Public  South... Bachel...
```

Manipulating the Data

- We can go deeper with this logic
 - "or" denoted with `|`
 - "and" denoted with `&`

```
df %>%  
  filter(str_detect(instnm, "Vand") & str_detect(instnm, "Univ"))
```

```
## # A tibble: 1 × 16  
##   unitid instnm  stabbr grad_debt_mdn control region preddeg  
##   <int> <chr>   <chr>         <int> <chr>   <chr> <chr>  
## 1 221999 Vander... TN             14962 Private South... Bachel...  
## # i 9 more variables: openadmp <int>, adm_rate <dbl>,  
## #   ccbasic <int>, sat_avg <int>, md_earn_wne_p6 <int>,  
## #   ugds <int>, costt4_a <int>, selective <dbl>,  
## #   research_u <dbl>
```

Manipulating the Data

- Can also put | or & in a single `str_detect()`

```
df %>%  
  filter(str_detect(instnm, 'Vand|Tenn'))
```

```
## # A tibble: 12 × 16  
##   unitid instnm stabbr grad_debt_mdn control region preddeg  
##   <int> <chr>   <chr>         <int> <chr>   <chr>   <chr>  
## 1 149639 Vande... IL           27000 Private Great... Bachel...  
## 2 220075 East ... TN           20500 Public  South... Bachel...  
## 3 220978 Middl... TN           21500 Public  South... Bachel...  
## 4 221485 South... TN              NA Public  South... Associ...  
## 5 221731 Tenne... TN           21500 Private South... Bachel...  
## 6 221740 The U... TN           20635 Public  South... Bachel...  
## 7 221759 The U... TN           20500 Public  South... Bachel...  
## 8 221768 The U... TN           22500 Public  South... Bachel...  
## 9 221838 Tenne... TN           27000 Public  South... Bachel...  
## 10 221847 Tenne... TN           17000 Public  South... Bachel...  
## 11 221999 Vande... TN           14962 Private South... Bachel...  
## 12 487010 The U... TN           13500 Public  South... Bachel...  
## # i 9 more variables: openadmp <int>, adm_rate <dbl>,  
## #   ccbasic <int>, sat avg <int>, md earn wne p6 <int>,
```


Manipulating the Data

- But **can't** do the same with `&`

```
df %>%  
  filter(str_detect(instnm, 'Vand&Univ'))
```

```
## # A tibble: 0 × 16  
## # i 16 variables: unitid <int>, instnm <chr>, stabbr <chr>,  
## #   grad_debt_mdn <int>, control <chr>, region <chr>,  
## #   preddeg <chr>, openadmp <int>, adm_rate <dbl>,  
## #   ccbasic <int>, sat_avg <int>, md_earn_wne_p6 <int>,  
## #   ugds <int>, costt4_a <int>, selective <dbl>,  
## #   research_u <dbl>
```

Manipulating the Data

- Negations are handled with `!`
 - Literally means "not"
- Drop rows with "of" in the school name

```
df %>%  
  filter(!str_detect(instnm, "of"))
```

```
## # A tibble: 2,025 × 16  
##   unitid instnm stabbr grad_debt_mdn control region preddeg  
##   <int> <chr>   <chr>         <int> <chr>   <chr>   <chr>  
## 1 100654 Alaba... AL             33375 Public  South... Bachel...  
## 2 100690 Amrid... AL             27334 Private South... Associ...  
## 3 100724 Alaba... AL             32000 Public  South... Bachel...  
## 4 100760 Centr... AL             12500 Public  South... Associ...  
## 5 100812 Athen... AL             19500 Public  South... Bachel...  
## 6 100830 Aubur... AL             24826 Public  South... Bachel...  
## 7 100858 Aubur... AL             21281 Public  South... Bachel...  
## 8 100937 Birmi... AL             25773 Private South... Bachel...  
## 9 101028 Chatt... AL             11931 Public  South... Associ...  
## 10 101161 Coast... AL             11000 Public  South... Associ...
```

Manipulating the Data

- (same as...)

```
df %>%  
  filter(!grepl("of",instnm))
```

```
## # A tibble: 2,025 × 16  
##   unitid instnm stabbr grad_debt_mdn control region preddeg  
##   <int> <chr>   <chr>         <int> <chr>   <chr>   <chr>  
## 1 100654 Alaba... AL             33375 Public  South... Bachel...  
## 2 100690 Amrid... AL             27334 Private South... Associ...  
## 3 100724 Alaba... AL             32000 Public  South... Bachel...  
## 4 100760 Centr... AL             12500 Public  South... Associ...  
## 5 100812 Athen... AL             19500 Public  South... Bachel...  
## 6 100830 Aubur... AL             24826 Public  South... Bachel...  
## 7 100858 Aubur... AL             21281 Public  South... Bachel...  
## 8 100937 Birmi... AL             25773 Private South... Bachel...  
## 9 101028 Chatt... AL             11931 Public  South... Associ...  
## 10 101161 Coast... AL             11000 Public  South... Associ...  
## # i 2,015 more rows  
## # i 9 more variables: openadmp <int>, adm_rate <dbl>,  
## #   ccbasic <int>, sat_avg <int>, md_earn_wne_p6 <int>,  
## #   ugds <int>, costt4 a <int>, selective <dbl>,
```

Manipulating: `select()`

- Still TMI!
- Before, I only cared about the admissions rate (`adm_rate`), the SAT scores (`sat_avg`), and the future earnings (`md_earn_wne_p6`)
- `select` will select **columns**

```
df %>%  
  filter(instnm == "Vanderbilt University") %>%  
  select(instnm, adm_rate, sat_avg, md_earn_wne_p6) # Select variables  
of interest
```

```
## # A tibble: 1 × 4  
##   instnm          adm_rate sat_avg md_earn_wne_p6  
##   <chr>          <dbl>   <int>         <int>  
## 1 Vanderbilt University  0.0912    1515         53400
```

Manipulating: `select()`

- We can use `matches()` function with `select()` in a manner similar to `str_detect()`

```
df %>%  
  select(matches("_"))
```

```
## # A tibble: 2,546 × 6  
##   grad_debt_mdn adm_rate sat_avg md_earn_wne_p6 costt4_a  
##   <int>      <dbl>   <int>      <int>      <int>  
## 1      33375    0.918     939      25200      23053  
## 2      22500    0.737    1234      35100      24495  
## 3      27334    NA         NA      30700      14800  
## 4      21607    0.826    1319      36200      23917  
## 5      32000    0.969     946      22600      21866  
## 6      23250    0.827    1261      37400      29872  
## 7      12500    NA         NA      23100      10493  
## 8      19500    NA         NA      33400         NA  
## 9      24826    0.904    1082      30100      19849  
## 10     21281    0.807    1300      39500      31590  
## # i 2,536 more rows  
## # i 1 more variable: research_u <dbl>
```

Stepping back

- **RQ**: How might admissions and SAT scores be **related**?
 - **Theory**: selective schools have stricter criteria
 - **Hypothesis**: admissions and SAT scores should be **negatively** related
- How can we test this hypothesis?

Summarizing Data: `summarise()` + `mean()`

- We can combine base R functions with `tidyverse` functions!
 - Base R: `mean()`
 - `tidyverse`: `summarise()` (aka `summarize()`)
- Overall average SAT scores

```
df %>%  
  summarise(mean_sat = mean(sat_avg, na.rm=T)) # Average SAT scores  
for entire data
```

```
## # A tibble: 1 × 1  
##   mean_sat  
##   <dbl>  
## 1    1141.
```

Summarizing Data

- Let's unpack this

```
df %>%  
  summarise(mean_sat = mean(sat_avg, na.rm=T))
```

- Create new variable `mean_sat` that contains the `mean()` of every school's average SAT score
- `na.rm=T` means we want to ignore missing data. If not?

```
df %>%  
  summarise(mean_sat = mean(sat_avg))
```

```
## # A tibble: 1 × 1  
##   mean_sat  
##   <dbl>  
## 1      NA
```


Summarizing Data

- Recall we want see if more selective schools have higher SAT scores

```
df %>%  
  filter(adm_rate < .1) %>% # Only schools who accept < 10%  
  summarise(mean_sat_LT10 = mean(sat_avg,na.rm=T)) # Average SAT
```

```
## # A tibble: 1 × 1  
##   mean_sat_LT10  
##           <dbl>  
## 1         1510.
```

```
df %>%  
  filter(adm_rate > .1) %>% # Only schools who accept > 10%  
  summarise(mean_sat_GT20 = mean(sat_avg,na.rm=T)) # Average SAT
```

```
## # A tibble: 1 × 1  
##   mean_sat_GT20  
##           <dbl>  
## 1         1135.
```

Adding / changing variables: `mutate()`

- `mutate()` creates a new variable

```
df %>%  
  mutate(newvar = 1) %>%  
  select(instnm, newvar)
```

```
## # A tibble: 2,546 × 2  
##   instnm                                newvar  
##   <chr>                                <dbl>  
## 1 Alabama A & M University             1  
## 2 University of Alabama at Birmingham  1  
## 3 Amridge University                   1  
## 4 University of Alabama in Huntsville  1  
## 5 Alabama State University             1  
## 6 The University of Alabama            1  
## 7 Central Alabama Community College    1  
## 8 Athens State University              1  
## 9 Auburn University at Montgomery      1  
## 10 Auburn University                   1
```

Object Assignment Operator: <-

- Thus far, nothing we have done has changed `df`
- Use object assignment operator `<-` to **overwrite** an existing object

```
df <- df %>%  
  mutate(adm_rate_pct = adm_rate*100)
```

- Did it work?

```
df %>%  
  summarise(adm_rate_pct = mean(adm_rate_pct,na.rm=T),  
            adm_rate = mean(adm_rate,na.rm=T))
```

```
## # A tibble: 1 × 2  
##   adm_rate_pct adm_rate  
##       <dbl>    <dbl>  
## 1       67.9    0.679
```

Logic: `ifelse()`

- 3 inputs:
 - Logical statement (labeled `test`)
 - Value if the logic is `TRUE` (labeled `yes`)
 - Value if the logic is `FALSE` (labeled `no`)
- `ifelse([LOGIC],[VALUE IF TRUE],[VALUE IF FALSE])`

Logic: `ifelse()`

- Say it out loud: "Create a new variable called `sel` that records if the school is selective or not. If the admissions rate is less than 10% (0.1), record the school as `sel = 1`. Otherwise, record the school as `sel = 0`."

```
df %>%  
  mutate(sel = ifelse(test = [LOGIC],  
                        yes = [VALUE IF TRUE],  
                        no = [VALUE IF FALSE]))
```

Logic: `ifelse()`

- Say it out loud: "Create a new variable called `sel` that records if the school is selective or not. **If the admissions rate is less than 10% (0.1)**, record the school as `sel = 1`. Otherwise, record the school as `sel = 0`."

```
df %>%  
  mutate(sel = ifelse(test = adm_rate < 0.1, # This is the logic  
                        yes = [VALUE IF TRUE],  
                        no = [VALUE IF FALSE]))
```

Logic: `ifelse()`

- Say it out loud: "Create a new variable called `sel` that records if the school is selective or not. If the admissions rate is less than 10% (0.1), **record the school as `sel = 1`**. Otherwise, record the school as `sel = 0`."

```
df %>%  
  mutate(sel = ifelse(test = adm_rate < 0.1, # This is the logic  
                      yes = 1, # This is the value if TRUE  
                      no = [VALUE IF FALSE]))
```

Logic: `ifelse()`

- Say it out loud: "Create a new variable called `sel` that records if the school is selective or not. If the admissions rate is less than 10% (0.1), record the school as `sel = 1`. **Otherwise, record the school as `sel = 0`.**"

```
df %>%  
  mutate(sel = ifelse(test = adm_rate < 0.1, # This is the logic  
                        yes = 1, # This is the value if TRUE  
                        no = 0)) # This is the value if FALSE
```


Logic: `ifelse()` + `mutate()`

- Remember that if we want to keep this, we need the **assignment operator** `<-`

```
df <- df %>%  
  mutate(sel = ifelse(test = adm_rate < 0.1, # This is the logic  
                      yes = 1, # This is the value if TRUE  
                      no = 0)) # This is the value if FALSE
```

Summarizing Data: `group_by()`

- One final `tidyverse` function: `group_by()`
- Let's use the newly created `selective` variable which is either 1 or 0

```
df %>%  
  select(instnm,selective,adm_rate)
```

```
## # A tibble: 2,546 × 3  
##   instnm                selective adm_rate  
##   <chr>                <dbl>   <dbl>  
## 1 Alabama A & M University      0     0.918  
## 2 University of Alabama at Birmingham      0     0.737  
## 3 Amridge University           NA      NA  
## 4 University of Alabama in Huntsville      0     0.826  
## 5 Alabama State University          0     0.969  
## 6 The University of Alabama          0     0.827  
## 7 Central Alabama Community College      NA      NA  
## 8 Athens State University          NA      NA  
## 9 Auburn University at Montgomery        0     0.904  
## 10 Auburn University             0     0.807  
## # i 2,536 more rows
```

Summarizing Data: `group_by()`

- Instead of running two separate `filter()` commands, use `group_by()`

```
df %>%  
  # Group the data by selective (either 1 or 0)  
  group_by(selective) %>%  
  # Calculate average SAT for each group  
  summarise(mean_sat = mean(sat_avg, na.rm=T))
```

```
## # A tibble: 3 × 2  
##   selective mean_sat  
##   <dbl>     <dbl>  
## 1      0     1135.  
## 2      1     1510.  
## 3     NA      NaN
```

Results

- Do more selective schools have higher SAT scores?
- Yes
- This **Result** **confirms** our **Hypothesis** and **answers** our **Research Question**

Conclusion

- What we've done today is a microcosm of data science
 1. Opened **data** (`readRDS`)
 2. Looked at **data** (`tidyverse` + `select()`, `filter()`, `arrange()`)
 3. Generated **hypotheses** (Admissions versus SAT scores)
 4. **Tested hypotheses** (`summarise()` + `mean()`)

Advanced Logic: `filter()`

If no time, jump to end

- `filter()` command with other logical operators
 - `>`, `<`: greater than, less than (`>=`, `<=`)
 - `!`: not (i.e., `!=` means "not equal to")
 - `&`: and
 - `|`: or

```
df %>%  
  # Schools EXCEPT Vandy  
  filter(instnm != "Vanderbilt University") %>%  
  select(instnm, stabbr, adm_rate, sat_avg)
```

```
## # A tibble: 2,545 × 4  
##   instnm          stabbr adm_rate sat_avg  
##   <chr>          <chr>    <dbl>   <int>  
## 1 Alabama A & M University AL        0.918     939  
## 2 University of Alabama at Birming... AL        0.737    1234  
## 3 Amridge University AL        NA        NA  
## 4 University of Alabama in Huntsvi... AL        0.826    1319  
## 5 Alabama State University AL        0.969     946  
## 6 The University of Alabama AL        0.827    1261
```

Advanced Logic: `str_detect()`

- `filter()` command with other functions
 - `str_detect([VAR],[PATTERN])`: detect a string
 - `grep1([PATTERN],[VAR])`: also detects a string

```
df %>%  
  filter(str_detect(instnm, "Vanderbilt")) %>%  
  select(instnm, stabbr, adm_rate, sat_avg)
```

```
## # A tibble: 1 × 4  
##   instnm          stabbr adm_rate sat_avg  
##   <chr>          <chr>    <dbl>   <int>  
## 1 Vanderbilt University TN      0.0912   1515
```

Advanced Logic: `str_detect()`

- String detection is case sensitive!

```
df %>%  
  filter(str_detect(instnm, "VAND")) %>%  
  select(instnm, stabbr, adm_rate, sat_avg)
```

```
## # A tibble: 0 × 4  
## #   i 4 variables: instnm <chr>, stabbr <chr>,  
## #     adm_rate <dbl>, sat_avg <int>
```

```
df %>%  
  filter(str_detect(instnm, "anderbil")) %>%  
  select(instnm, stabbr, adm_rate, sat_avg)
```

```
## # A tibble: 1 × 4  
##   instnm          stabbr adm_rate sat_avg  
##   <chr>          <chr>    <dbl>   <int>  
## 1 Vanderbilt University TN      0.0912   1515
```


Advanced Logic: & (and), | (or)

```
df %>%  
  filter(str_detect(instnm, "Colorado")) %>%  
  select(instnm, stabbr, adm_rate, sat_avg)
```

```
## # A tibble: 12 × 4  
##   instnm                stabbr adm_rate sat_avg  
##   <chr>                <chr>    <dbl>   <int>  
## 1 University of Colorado Denver/An... CO      0.673    1124  
## 2 University of Colorado Colorado ... CO      0.872    1136  
## 3 University of Colorado Boulder      CO      0.784    1276  
## 4 Colorado Christian University      CO      NA        NA  
## 5 Colorado College                   CO      0.135     NA  
## 6 Colorado School of Mines           CO      0.531    1342  
## 7 Colorado State University-Fort C... CO      0.814    1204  
## 8 Colorado Mesa University           CO      0.782    1063  
## 9 University of Northern Colorado    CO      0.908    1096  
## 10 Colorado State University Pueblo   CO      0.930    1047  
## 11 Western Colorado University        CO      0.842    1114  
## 12 Colorado State University-Global... CO      0.986    1048
```

Advanced Logic: & (and), | (or)

```
df %>%  
  filter(grepl("Colorado",instnm) & grepl(' of ',instnm)) %>%  
  select(instnm,stabbr,adm_rate,sat_avg)
```

```
## # A tibble: 5 × 4  
##   instnm                stabbr adm_rate sat_avg  
##   <chr>                <chr>    <dbl>   <int>  
## 1 University of Colorado Denver/Ans... CO      0.673    1124  
## 2 University of Colorado Colorado S... CO      0.872    1136  
## 3 University of Colorado Boulder      CO      0.784    1276  
## 4 Colorado School of Mines            CO      0.531    1342  
## 5 University of Northern Colorado     CO      0.908    1096
```

Advanced Logic: & (and), | (or)

```
df %>%  
  filter(grepl("Colorado",instnm) | grepl('Vermont',instnm)) %>%  
  select(instnm,stabbr,adm_rate,sat_avg)
```

```
## # A tibble: 16 × 4  
##   instnm                stabbr adm_rate sat_avg  
##   <chr>                <chr>    <dbl>   <int>  
## 1 University of Colorado Denver/An... CO      0.673    1124  
## 2 University of Colorado Colorado ... CO      0.872    1136  
## 3 University of Colorado Boulder     CO      0.784    1276  
## 4 Colorado Christian University      CO      NA        NA  
## 5 Colorado College                   CO      0.135     NA  
## 6 Colorado School of Mines           CO      0.531    1342  
## 7 Colorado State University-Fort C... CO      0.814    1204  
## 8 Colorado Mesa University           CO      0.782    1063  
## 9 University of Northern Colorado    CO      0.908    1096  
## 10 Colorado State University Pueblo   CO      0.930    1047  
## 11 Western Colorado University        CO      0.842    1114  
## 12 Community College of Vermont       VT      NA        NA  
## 13 Northern Vermont University        VT      0.778     NA  
## 14 Vermont Technical College          VT      0.670     NA  
## 15 University of Vermont              VT      0.673    1287
```

Advanced Logic: & (and), | (or)

```
df %>%  
  filter((grepl("Colorado",instnm) | grepl('Vermont',instnm)) &  
grepl(' of ',instnm)) %>%  
  select(instnm,stabbr,adm_rate,sat_avg)
```

```
## # A tibble: 7 × 4  
##   instnm          stabbr adm_rate sat_avg  
##   <chr>          <chr>    <dbl>   <int>  
## 1 University of Colorado Denver/Ans... CO        0.673    1124  
## 2 University of Colorado Colorado S... CO        0.872    1136  
## 3 University of Colorado Boulder      CO        0.784    1276  
## 4 Colorado School of Mines            CO        0.531    1342  
## 5 University of Northern Colorado     CO        0.908    1096  
## 6 Community College of Vermont       VT         NA       NA  
## 7 University of Vermont              VT        0.673    1287
```

Advanced Logic: & (and), | (or)

- & can be separated into multiple `filter()` commands

```
df %>%  
  filter((grepl("Colorado",instnm) | grepl('Vermont',instnm))) %>%  
  filter(grepl(' of ',instnm)) %>%  
  select(instnm,stabbr,adm_rate,sat_avg)
```

```
## # A tibble: 7 × 4  
##   instnm                stabbr adm_rate sat_avg  
##   <chr>                <chr>    <dbl>   <int>  
## 1 University of Colorado Denver/Ans... CO      0.673    1124  
## 2 University of Colorado Colorado S... CO      0.872    1136  
## 3 University of Colorado Boulder      CO      0.784    1276  
## 4 Colorado School of Mines            CO      0.531    1342  
## 5 University of Northern Colorado      CO      0.908    1096  
## 6 Community College of Vermont        VT      NA      NA  
## 7 University of Vermont               VT      0.673    1287
```

Advanced Logic: & (and), | (or)

- | can be moved into the `str_detect()` or `grepl()` commands

```
df %>%  
  filter(grepl("Colorado|Vermont", instnm)) %>%  
  filter(grepl(' of ', instnm)) %>%  
  select(instnm, stabbr, adm_rate, sat_avg)
```

```
## # A tibble: 7 × 4  
##   instnm                stabbr adm_rate sat_avg  
##   <chr>                <chr>    <dbl>   <int>  
## 1 University of Colorado Denver/Ans... CO      0.673    1124  
## 2 University of Colorado Colorado S... CO      0.872    1136  
## 3 University of Colorado Boulder      CO      0.784    1276  
## 4 Colorado School of Mines            CO      0.531    1342  
## 5 University of Northern Colorado     CO      0.908    1096  
## 6 Community College of Vermont        VT      NA        NA  
## 7 University of Vermont               VT      0.673    1287
```

Quick Test

- Filter schools from Texas with the word "community" in their name

```
# INSERT CODE HERE
```

Advanced Logic: `select()`

- `select` can be paired with `matches()` or `contains()` for similar flexibility (equivalent to `str_detect()` or `grep1()` for `filter()`)

```
df %>%  
  select(contains('inst'))
```

```
## # A tibble: 2,546 × 1  
##   instnm  
##   <chr>  
## 1 Alabama A & M University  
## 2 University of Alabama at Birmingham  
## 3 Amridge University  
## 4 University of Alabama in Huntsville  
## 5 Alabama State University  
## 6 The University of Alabama  
## 7 Central Alabama Community College  
## 8 Athens State University  
## 9 Auburn University at Montgomery  
## 10 Auburn University  
## # i 2,536 more rows
```


Advanced Logic: `select()`

- `matches` can work with `|`

```
df %>%  
  select(!matches('_|inst'))
```

```
## # A tibble: 2,546 × 10  
##   unitid stabbr control region    preddeg openadmp ccbasic  
##   <int> <chr>   <chr>   <chr>    <chr>      <int>   <int>  
## 1 100654 AL      Public Southeast Bachelo...      2     18  
## 2 100663 AL      Public Southeast Bachelo...      2     15  
## 3 100690 AL      Private Southeast Associa...      1     20  
## 4 100706 AL      Public Southeast Bachelo...      2     16  
## 5 100724 AL      Public Southeast Bachelo...      2     19  
## 6 100751 AL      Public Southeast Bachelo...      2     15  
## 7 100760 AL      Public Southeast Associa...      1      2  
## 8 100812 AL      Public Southeast Bachelo...     NA     22  
## 9 100830 AL      Public Southeast Bachelo...      2     18  
## 10 100858 AL      Public Southeast Bachelo...      2     15  
## # i 2,536 more rows  
## # i 3 more variables: ugds <int>, selective <dbl>,  
## #   sel <dbl>
```

Advanced Logic: `select()`

- `select` can also work with `where` to find classes

```
df %>%  
  select(where(is.numeric))
```

```
## # A tibble: 2,546 × 13  
##   unitid grad_debt_mdn openadmp adm_rate ccbasic sat_avg  
##   <int>      <int>      <int>    <dbl>    <int>    <int>  
## 1 100654      33375         2    0.918      18      939  
## 2 100663      22500         2    0.737      15     1234  
## 3 100690      27334         1    NA         20      NA  
## 4 100706      21607         2    0.826      16     1319  
## 5 100724      32000         2    0.969      19      946  
## 6 100751      23250         2    0.827      15     1261  
## 7 100760      12500         1    NA         2      NA  
## 8 100812      19500        NA    NA         22      NA  
## 9 100830      24826         2    0.904      18     1082  
## 10 100858      21281         2    0.807      15     1300  
## # i 2,536 more rows  
## # i 7 more variables: md_earn_wne_p6 <int>, ugds <int>,  
## #   costt4_a <int>, selective <dbl>, research_u <dbl>,  
## #   adm rate pct <dbl>, sel <dbl>
```

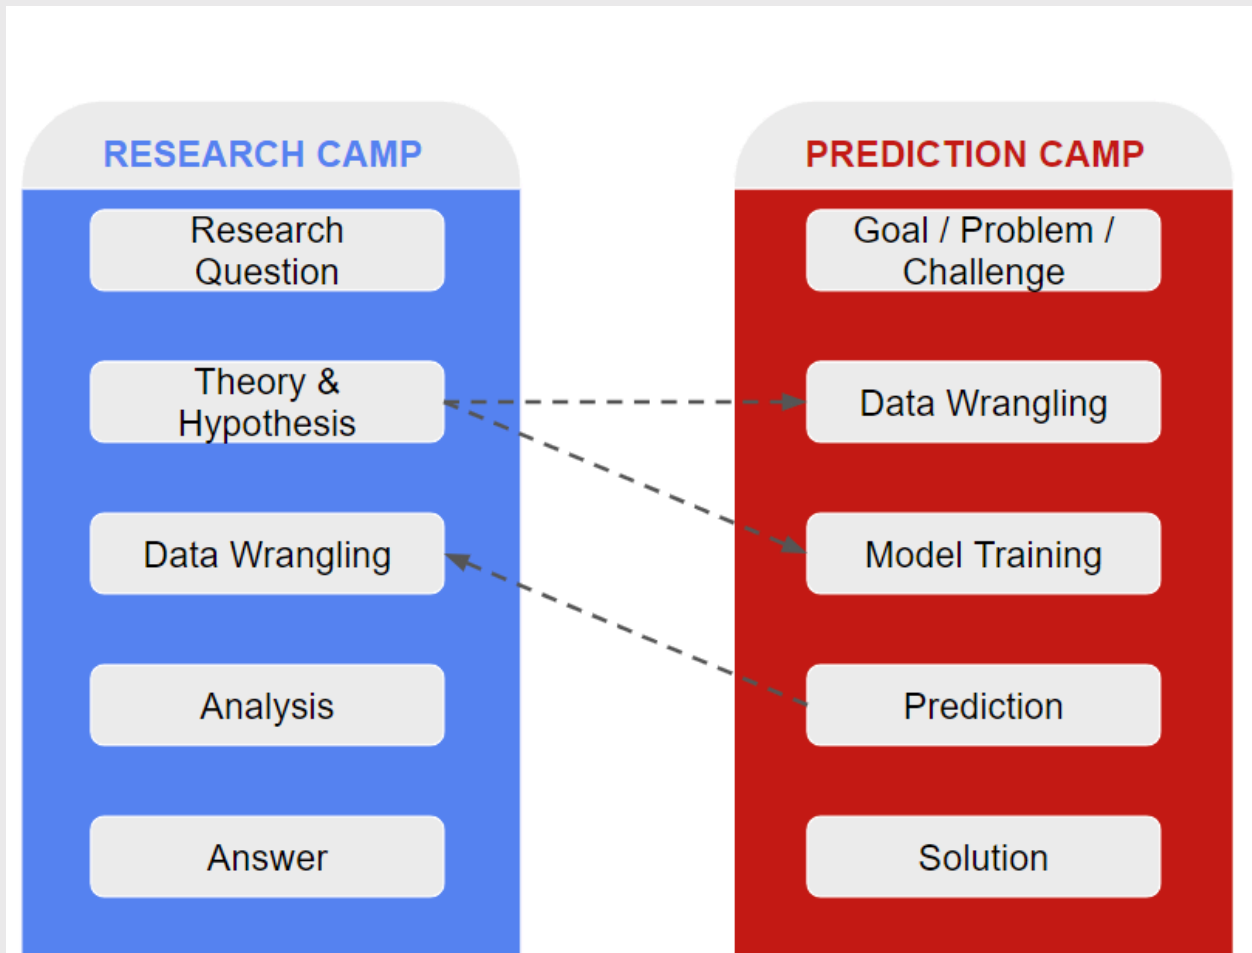
Quick Test

- Filter to only schools in California and select only character columns

```
# INSERT CODE HERE
```

BREAK

The Two Camps



The Research Camp

- RQ: How might admissions and SAT scores be **related**?
 - Theory: selective schools have stricter criteria
 - Hypothesis: admissions and SAT scores should be **negatively** related
- How can we test this hypothesis?

Previously: `summarise()`

- We can combine base R functions with `tidyverse` functions!
 - Base R: `mean()`
 - `tidyverse`: `summarise()` (aka `summarize()`)
- Overall average SAT scores

```
df %>%  
  summarise(mean_sat = mean(sat_avg, na.rm=T))
```

```
## # A tibble: 1 × 1  
##   mean_sat  
##   <dbl>  
## 1    1141.
```


Previously: `summarise()`

- Let's unpack this

```
df %>%  
  summarise(mean_sat = mean(sat_avg, na.rm=T))
```

- Create new variable `mean_sat` that contains the `mean()` of every school's average SAT score
- `na.rm=T` means we want to ignore missing data. If not?

```
df %>%  
  summarise(mean_sat = mean(sat_avg))
```

```
## # A tibble: 1 × 1  
##   mean_sat  
##   <dbl>  
## 1      NA
```

summarise() + filter()

- Recall we want see if more selective schools have higher SAT scores

```
df %>%  
  filter(adm_rate < .1) %>%  
  summarise(mean_sat_LT10 = mean(sat_avg, na.rm=T))
```

```
## # A tibble: 1 × 1  
##   mean_sat_LT10  
##           <dbl>  
## 1          1510.
```

```
df %>%  
  filter(adm_rate > .1 & adm_rate < .2) %>%  
  summarise(mean_sat_1020 = mean(sat_avg, na.rm=T))
```

```
## # A tibble: 1 × 1  
##   mean_sat_1020  
##           <dbl>  
## 1          1424.
```

summarise() + group_by()

- One final `tidyverse` function: `group_by()`

```
df %>%  
  group_by(selective) %>%  
  summarise(mean_sat = mean(sat_avg, na.rm=T))
```

```
## # A tibble: 3 × 2  
##   selective mean_sat  
##   <dbl>     <dbl>  
## 1      0     1135.  
## 2      1     1510.  
## 3     NA      NaN
```


Plotting data

- Let's plot the data instead of writing many of these `summarise()` functions
- Visualization in R uses `ggplot()` function
 - Inputs: `aes(x,y,...)` (elipses `...` indicates many more inputs)
 - `x` is the x-axis (horizontal)
 - `y` is the y-axis (vertical)

ggplot()

- Attach `ggplot()` to your data with `%>%`

```
df %>%  
  ggplot()
```




ggplot()

- Then tell it what to put in the x-axis and y-axis
- What should go on these axes?
- **Theory**: Selective schools choose higher scoring students
 - Selective schools **explain** higher scores
 - Selective schools: **independent variable** / **explanatory variable** / **predictor** / X
 - Higher scores: **dependent variable** / **outcome variable** / Y
- Selective schools go on the x-axis, SAT scores go on the y-axis

ggplot()

```
df %>%  
  ggplot(aes(x = adm_rate, y = sat_avg))
```



ggplot()


- This gives us an empty plot
- We have the correct variables on the correct axes...
- ...but we need to choose how to display them
- There are many different `ggplot()` functions to choose from
 - `geom_point()` creates one point for each x and y coordinate
 - `geom_bar()` creates a barplot
 - `geom_histogram()` creates a histogram
 - `geom_density()` creates a density plot
 - `geom_boxplot()` creates a box-and-whisker plot

ggplot()

- We **add** a second `ggplot()` function to the first with a plus sign +
 - **NB:** This is JUST LIKE THE PIPE OPERATOR `%>%` in `tidyverse`!
- Since `adm_rate` (the x-axis variable) and `sat_avg` (the y-axis variable) are both numeric ("continuous") measures, we will use `geom_point()`
 - We will come back to **variable types** and how to visualize them later

ggplot()

```
df %>%  
  ggplot(aes(x = adm_rate, y = sat_avg)) +  
  geom_point()
```




Plotting data

- Let's unpack this
 - `aes(x,y)` sets the basic aesthetics for the plot
 - `geom_point()` tells `ggplot()` how to visualize those aesthetics
 - These two parts are linked with the `+`. Similar to...?
 - ...the `%>%` in `tidyverse`!

Interpreting the plot

- We **hypothesized** that admissions and SAT scores are negatively related
 - Is this supported in the data?
- Let's add a line of best fit with `geom_smooth()`

```
df %>%  
  ggplot(aes(x = adm_rate, y = sat_avg)) +  
  geom_point() +  
  geom_smooth(method = 'lm', se = F)
```



The Research Camp

- **RQ:** How might future earnings and SAT scores be **related**?
 - **Theory:** SATs measure student ability.
 - **Theory:** Student ability is valued by the labor market.
 - **Theory:** Firms pay more for students with higher SAT scores.
 - **Hypothesis:** Earnings and SAT scores should be **positively** related

Plotting Quiz

- Which variable goes on the x-axis?
 - **SAT scores**
- Which variable goes on the y-axis?
 - **Earnings**
- In our theory, SAT scores **cause** earnings
- Why might this **not** be the case?
 - Spurious 1: SAT scores **and** earnings are caused by student ability
 - Spurious 2: SAT scores **and** earnings are caused by socio-economic privilege

Let's Plot!

```
df %>%  
  ggplot(aes(x = sat_avg, y = md_earn_wne_p6)) + # Build axes  
  geom_point() + # Add points  
  geom_smooth(method = 'lm', se = F) # Add line of best fit
```


Outliers

- Which schools are furthest from the line?
 - These are **outliers**
 - These schools are the **furthest** from our [theory](#)

```
df %>%  
  mutate(out = ifelse(md_earn_wne_p6 > 100000,  
                      instnm, # Value if TRUE  
                      NA)) %>% # Value if FALSE  
  drop_na(out,sat_avg) %>%  
  select(instnm,md_earn_wne_p6,sat_avg)
```


```
## # A tibble: 2 × 3  
##   instnm                md_earn_wne_p6 sat_avg  
##   <chr>                <int>     <int>  
## 1 University of Health Sciences and ... 120400    1262  
## 2 Albany College of Pharmacy and Hea... 112100    1242
```

Plotting data

- We can add these as labels!

```
df %>%  
  mutate(out = ifelse(md_earn_wne_p6 > 100000,  
                      instnm, # Value if TRUE  
                      NA)) %>% # Value if FALSE  
  ggplot(aes(x = sat_avg, y = md_earn_wne_p6,  
             label = out)) +  
  geom_point() +  
  geom_text()
```

Plotting data




Plotting data

- Let's accentuate the outlier more with color

```
p <- df %>%  
  mutate(out = ifelse(md_earn_wne_p6 > 100000,  
                      instnm, # Value if TRUE  
                      NA)) %>% # Value if FALSE  
  
  drop_na(sat_avg) %>%  
  ggplot(aes(x = sat_avg, y = md_earn_wne_p6,  
            label = out, color = out)) +  
  geom_point() +  
  scale_color_manual(name = "Outlier", values =  
c('red', 'red', 'black')) +  
  geom_text(hjust = .5, vjust = 1, color = 'black', size = 3)
```

Plotting data



Categorical Data

- Thus far, plotting two continuous variables with `geom_point()`
- What if we wanted to see which state has the most selective schools?
- Use `group_by()` and `summarise()`

```
df %>%  
  group_by(stabbr) %>%  
  summarise(selective_avg = mean(adm_rate, na.rm=T))
```

```
## # A tibble: 51 × 2  
##   stabbr selective_avg  
##   <chr>         <dbl>  
## 1 AK           0.827  
## 2 AL           0.654  
## 3 AR           0.676  
## 4 AZ           0.843  
## 5 CA           0.592  
## 6 CO           0.768  
## 7 CT           0.589  
## 8 DC           0.529  
## 9 DE           0.627
```

Categorical Data


- Two variables (`stabbr` and `selective_avg`), but one of them is now a `character` type
- Can we plot this as a scatterplot?

```
p <- df %>%  
  group_by(stabbr) %>%  
  summarise(selective_avg = mean(adm_rate, na.rm=T)) %>%  
  ggplot(aes(x = stabbr, y = selective_avg)) +  
  geom_point()
```

Categorical Data

- Yes...but it isn't very pretty


p



Categorical Data: `geom_bar()`

- NB: `geom_bar()` will automatically count the values on the x-axis

```
df %>%  
  ggplot(aes(x = stabbr)) +  
  geom_bar()
```




Categorical Data: `geom_bar()`

- This is fine if we just want to know which states have the most schools in our data
- But we want to put the average admissions rate on the y-axis instead
 - Need to **override** `geom_bar()` default behavior

```
p <- df %>%  
  group_by(stabbr) %>%  
  summarise(selective_avg = mean(adm_rate, na.rm=T)) %>%  
  ggplot(aes(x = stabbr, y = selective_avg)) +  
  geom_bar(stat = 'identity')
```

Categorical Data

p



Categorical Data


- Getting a little better, but still ugly
- Use `reorder()` to sort the x-axis values by the y-axis

```
p <- df %>%  
  group_by(stabbr) %>%  
  summarise(selective_avg = mean(adm_rate, na.rm=T)) %>%  
  ggplot(aes(x = reorder(stabbr, selective_avg), y = selective_avg)) +  
  geom_bar(stat = 'identity')
```

Categorical Data

- Even better!

p




Plot Tweaking

- We could go even further and swap the x and y-axes (although this isn't always a good idea!)

```
p <- df %>%  
  group_by(stabbr) %>%  
  summarise(selective_avg = mean(adm_rate,na.rm=T)) %>%  
  ggplot(aes(y = reorder(stabbr,selective_avg),x = selective_avg)) +  
  geom_bar(stat = 'identity')
```

Plot Tweaking


p



- Still ugly though! We want to tweak the labels with `labs()`

Plot Tweaking

```
p +  
  labs(title = "Average Admissions by State",  
        x = "Admissions Rate",  
        y = "State")
```



Conclusion

- What to take away
 1. Which variables go on which axes
 2. How to put these on a `ggplot()` figure
 3. How to create a visualization of these variables
- This wraps up the crash course in R
 - **REMEMBER:** This class is *inherently* challenging because of R
 - The course is graded leniently to reflect the inherent difficulty of the material