

Regression

Part 1

Prof. Bisbee

Seoul National University

Slides Updated: 2024-07-16

Agenda


1. Modeling Conditional Variation
2. Adding Regression to the **Process**
3. Introducing the Data
4. Demonstrating Regressions

Regression & Conditional Analysis

- Recall our discussion of **conditional** analysis
 - Conditional → **depends on**
 - Analyze with **conditional means**


What is regression?

- Conditional means for continuous data




What is regression?

- People with hourly wages < \$20 spend ~\$50 on entertainment per week




What is regression?

- People with hourly wages $> \$40$ spend $\sim \$95$ on entertainment per week




What is regression?

- **Theory**: the more you earn, the more you spend




What is regression?

- But **conditional means** make a lot of mistakes. Can we do better?



What is regression?


- But **conditional means** make a lot of mistakes. Can we do better?




Regression

- Calculating a **line** that minimizes mistakes *for every observation*
 - NB: could be a curvey line! For now, just assume straight
- Recall from geometry how to graph a straight line
- $Y = a + bX$
 - a : the "intercept" (where the line intercepts the y-axis)
 - b : the "slope" (how much Y changes for each increase in X)
- (Data scientists use α and β instead of a and b b/c nerds)
- Regression analysis simply chooses the best line
 - "Best"?
 - The line that minimizes the mistakes (the **line of best fit**)


Linear Regression



Linear Regression




Linear Regression




- **Error/Residual:** mistake made by a line
 - In math: $u_i = y_i - \hat{y}_i$
 - In English: difference between true outcome value (y_i) and prediction (\hat{y}_i)

Linear Regression




- Use **errors** to find **line of best fit**
- **RMSE** (**R**oot **M**ean **S**quared **E**rror)
 - Square the errors
 - Take their average
 - Take the square root
- **RMSE** = 1.23

Linear Regression




- Use **errors** to find **line of best fit**
- **RMSE (Root Mean Squared Error)**
 - Square the errors
 - Take their average
 - Take the square root
- **RMSE = 1.48**

Linear Regression




- Use **errors** to find **line of best fit**
- **RMSE** (**R**oot **M**ean **S**quared **E**rror)
 - Square the errors
 - Take their average
 - Take the square root
- **RMSE** = 2.19

Linear Regression




- Use **errors** to find **line of best fit**
- **RMSE** (**R**oot **M**ean **S**quared **E**rror)
 - Square the errors
 - Take their average
 - Take the square root
- **RMSE** = 1.03

Visual Intuition



Visual Intuition



Regression

- The line is **substantively meaningful**
- Red line on scatter plot of spending and wages: $Y = \underbrace{12}_{\alpha} + \underbrace{2}_{\beta} * X$
- α tells us the value of Y when X is zero
 - People who don't make any money spend \$12 per week on entertainment
- β tells us how much Y increases for each additional X
 - People spend an additional \$2 per week for each additional \$1 in hourly wages

Regression

- These are called "**linear models**"
 - **Not** because the line is straight (it might not be)
 - but because the components are additive ($\alpha + \beta X$)
- Can extend to multiple predictors (X 's)
 - $Y = \alpha + \beta_1 X_1 + \beta_2 X_2 + \dots + \varepsilon$
 - X_1 might be wages and X_2 might be age (for example)
 - The final term ε measures how bad our mistakes are

Regression

- Let's demonstrate with the `debt` data

```
require(tidyverse)

debt <-
read_rds('https://github.com/jbisbee1/ISP_Data_Science_2024/raw/main/debt.rds')

glimpse(debt)
```

```
## Rows: 2,546
## Columns: 16
## $ unitid      <int> 100654, 100663, 100690, 100706, 100...
## $ instnm      <chr> "Alabama A & M University", "Univer...
## $ stabbr      <chr> "AL", "AL", "AL", "AL", "AL", "AL", ...
## $ grad_debt_mdn <int> 33375, 22500, 27334, 21607, 32000, ...
## $ control     <chr> "Public", "Public", "Private", "Pub...
## $ region      <chr> "Southeast", "Southeast", "Southeas...
## $ preddeg     <chr> "Bachelor's", "Bachelor's", "Associ...
## $ openadmp     <int> 2, 2, 1, 2, 2, 2, 1, NA, 2, 2, 2, 1...
## $ adm_rate     <dbl> 0.9175, 0.7366, NA, 0.8257, 0.9690, ...
## $ ccbasic      <int> 18, 15, 20, 16, 19, 15, 2, 22, 18, ...
## $ sat_avg      <int> 939, 1234, NA, 1319, 946, 1261, NA, ...
```

Research Camp

- **Research Question:** What is the relationship between SAT scores and median future earnings?
- **Theory:** Students with higher SAT scores work harder and have learned more. Employers reward these attributes with higher wages in the private market.
- **Hypothesis:** The relationship between SAT scores and future earnings should be positive.
 - **NB:** Important caveats to this simplistic theory!
 - Socioeconomic status: predicts both higher SAT scores and higher wages
 - **Correlation \neq Causation**

Set Up

- Linking Theory to Data
- Our SAT scores are theorized to explain future earnings
 - Thus the SAT scores are the independent / explanatory / predictor variable X
 - And earnings are the dependent / outcome variable Y

Regression

- There is a simple recipe to follow
- And it is exactly how the syllabus for the class is designed!
 1. Look at your data to identify missingness (**Wrangling: Lecture 3**)
 2. **Univariate** visualization of your variables (**Lecture 4**)
 3. **Multivariate** visualization of your variables (**Lecture 5**)
 4. **Regression** (today)
 5. Evaluation of **errors** (today)

Step 1: Look

- Why worry about **missingness**?
1. **Substantive:** external validity
 2. **Technical:** cross validation won't work!

```
summary(debt %>% select(sat_avg,md_earn_wne_p6))
```

```
##      sat_avg      md_earn_wne_p6
## Min.       : 737   Min.       : 10600
## 1st Qu.:1053   1st Qu.: 26100
## Median :1119   Median : 31500
## Mean    :1141   Mean    : 33028
## 3rd Qu.:1205   3rd Qu.: 37400
## Max.    :1557   Max.    :120400
## NA's    :1317   NA's    :240
```


Step 2: Univariate Viz

- Why visualize both Y and X ?
 1. **Substantive:** See which units you are talking about
 2. **Technical:** Adjust for *skew*

Step 2: Univariate Viz

- Why visualize both Y and X ?


```
debt %>%  
  ggplot(aes(x = sat_avg)) +  
  geom_histogram()
```



Step 2: Univariate Viz

- Why visualize both Y and X ?


```
debt %>%  
  ggplot(aes(x = md_earn_wne_p6)) +  
  geom_histogram()
```



Step 3: Multivariate

- Eyeball the relationship first!


```
debt %>%  
  ggplot(aes(x = sat_avg, y = md_earn_wne_p6)) +  
  geom_point()
```



Step 3: Multivariate Viz

- Adding regression line

```
debt %>%  
  ggplot(aes(x = sat_avg, y = md_earn_wne_p6)) +  
  geom_point() + geom_smooth(method = 'lm', se = F)
```



Step 3: Multivariate Viz


- Let's focus on two schools

```
toplot <- debt %>%  
  mutate(hl = ifelse(unitid %in% c(100654,179265), 'hl', 'none')) #  
Choosing two examples  
p2 <- toplot %>%  
  ggplot(aes(x = sat_avg, y = md_earn_wne_p6,color = hl,group =  
1,alpha = hl)) +  
  geom_point(data = toplot %>% filter(hl == 'none')) +  
  geom_point(data = toplot %>% filter(hl == 'hl'),size =3) +  
  scale_alpha_manual(values = c(1,.3)) +  
  scale_color_manual(values = c('red','black')) +  
  geom_smooth(method = 'lm',se = F) +  
  theme(legend.position = 'none') +  
  labs(title = "Graduate Earnings and SAT Scores",  
        subtitle = "By School",  
        x = "Average SAT Score",  
        y = "Median Earnings of Grads")
```


Step 3: Multivariate Viz

- Adding regression line

p2



Step 3: Multivariate Viz


- Defining ε

```
p3 <- toplot %>%  
  ggplot(aes(x = sat_avg, y = md_earn_wne_p6,color = h1,group =  
1,alpha = h1)) +  
  geom_point(data = toplot %>% filter(h1 == 'none')) +  
  geom_point(data = toplot %>% filter(h1 == 'h1'),size =3) +  
  scale_alpha_manual(values = c(1,.3)) +  
  scale_color_manual(values = c('red','black')) +  
  geom_smooth(method = 'lm',se = F) +  
  annotate(geom = 'segment',  
    x = toplot %>% filter(h1 == 'h1') %>% .$sat_avg,  
    y = toplot %>% filter(h1 == 'h1') %>% .$md_earn_wne_p6,  
    xend = toplot %>% filter(h1 == 'h1') %>% .$sat_avg,  
    yend = c(27500,41000),color = 'red',lwd = 1.2) +  
  theme(legend.position = 'none') +  
  labs(title = "Graduate Earnings and SAT Scores",  
    subtitle = "By School",  
    x = "Average SAT Score",  
    y = "Median Earnings of Grads")
```

Step 3: Multivariate Viz

- Measuring errors

p3




The Data Scientist's Trade-off

- Those mistakes seem pretty big!
- Why not use a curvier line?


The Data Scientist's Trade-off

- Those mistakes seem pretty big!
- Why not use a curvier line?




The Data Scientist's Trade-off

- Those mistakes seem pretty big!
- Why not use a curvier line?




The Data Scientist's Trade-off

- Those mistakes seem pretty big!
- Why not use a curvier line?



The Data Scientist's Trade-off

- Those mistakes seem pretty big!
- Why not use a curvier line?



The Data Scientist's Trade-off

- Want to **reduce complexity**
- But also want to be **accurate**
- What is the right answer?
 - It depends on your **theory** and the **data**
 - It is context-dependent
- And this is still only using *linear regression models*!
 - This is a deep area of study, for those interested

Step 4: Regression

- Introducing the `lm(formula,data)` function
- Two inputs to care about:
 - `formula`: Code for $Y = \alpha + \beta X$
 - `data`: What is the data we are using?
- `formula` is written as $Y \sim X$
 - R will calculate α and β for us
 - Just need to tell it what is Y (`md_earn_wne_p6`) and X (`sat_avg`)
 - The tilde (`~`) is R's version of the equals sign **in a regression equation**
- Save the model to an object

```
model_earn_sat <- lm(formula = md_earn_wne_p6 ~ sat_avg, data = debt)
```

Step 4: Interpretation

- What is in this object?
- The regression results! Look at them with `summary()`

```
summary(model_earn_sat)
```

```
##
## Call:
## lm(formula = md_earn_wne_p6 ~ sat_avg, data = debt)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -23239  -4311   -852    2893   78695
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -12053.87    1939.80  -6.214 7.12e-10 ***
## sat_avg      42.60       1.69   25.203 < 2e-16 ***
## ---
## Signif. codes:
## 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

Step 4: Interpretation

- Starting with the first column called **Estimate**
- 1st Row (**Intercept**) is α : the predicted value of Y when X is zero
 - Schools with average SAT scores of 0 produce graduates who earn -\$12,053.87
 - Sensible?
- 2nd Row **sat_avg** is the β : the increase in Y when X increases by one
 - For each unit increase in the average SAT score, recent graduates earn \$42.60 more
 - Sensible?

Step 4: Interpretation

- Other 3 columns?
 - `Std. Error` is the "standard error"
 - `t value` is the "t-statistic"
 - `Pr(>|t|)` is the "p-value"
- $t\text{-statistic} = \text{Estimate} / \text{standard error}$
- $p\text{-value} = \text{function}(t\text{-statistic})$
 - Only really need to remember the p-value for this course
 - This is 1 minus confidence
 - The lower the p-value, the **more** confident we are that the `Estimate` is **not zero**

Step 4: Interpretation

```
summary(model_earn_sat)
```

```
##
## Call:
## lm(formula = md_earn_wne_p6 ~ sat_avg, data = debt)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -23239  -4311   -852    2893   78695
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -12053.87    1939.80  -6.214 7.12e-10 ***
## sat_avg      42.60       1.69   25.203 < 2e-16 ***
## ---
## Signif. codes:
## 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7594 on 1196 degrees of freedom
## (1348 observations deleted due to missingness)
## Multiple R-squared:  0.3469,    Adjusted R-squared:  0.3463
## F-statistic: 635.2 on 1 and 1196 DF,  p-value: < 2.2e-16
```

Step 4: Interpretation

- Kinda ugly? Use `tidy()` function from the `broom` package

```
require(broom)
```

```
## Loading required package: broom
```

```
tidy(model_earn_sat)
```

```
## # A tibble: 2 × 5
##   term          estimate std.error statistic    p.value
##   <chr>          <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept) -12054.    1940.    -6.21 7.12e- 10
## 2 sat_avg      42.6      1.69     25.2 9.09e-113
```

Another Example

- We will come back to the RMSE after the break
- For now, let's try with a different research question!
- What is the relationship between admissions and future earnings?
 - Theory: More selective schools are more prestigious
 - Hypothesis: There should be a negative relationship between the admissions rate and future earnings

Do It Together!

1. Look at the data and acknowledge missingness
2. Univariate visualization of X and Y
3. Multivariate visualization of X and Y
4. Regression

BREAK

Learning Goals

1. Skew, logs, and coefficients
2. Evaluating a regression: Univariate and multivariate visualization of errors
3. Root Mean Squared Error (RMSE)
4. Cross Validation

Evaluating Regression Results

- Understanding the **errors** helps us evaluate the model
 - Define the errors $\varepsilon = Y - \hat{Y}$
 - True outcome values Y
 - Predicted outcome values \hat{Y}
 - Useful to assess model performance
1. **Look** with univariate and multivariate visualization of the errors
 2. Calculate the **RMSE**

Introducing the Data

- New dataset on **movies**
 - `require tidyverse`, and `plotly` packages
 - Load `mv.Rds` from Github to object `mv`

```
require(tidyverse)
```

```
mv <-
```

```
read_rds('https://github.com/jbisbee1/ISP_Data_Science_2024/raw/main/ds')
```

RQ: Hollywood Finances

- **Research Question:** What is the relationship between a movie's budget (how much it costs to make a movie) and a movie's earnings (how much money people pay to see the movie in theaters)?
- **Theory:** More money spent means more famous actors, better special effects, stronger marketing
- **Hypothesis:** earnings (**gross**) and costs (**budget**) should be **positively** correlated
 - X : ?
 - Y : ?

Follow the process: Look

- TONS of missingness!

```
summary(mv %>% select(gross,budget))
```

```
##           gross           budget
## Min.      :7.140e+02   Min.      :    5172
## 1st Qu.:1.121e+07   1st Qu.: 16865322
## Median :5.178e+07   Median : 37212044
## Mean    :1.402e+08   Mean    : 57420173
## 3rd Qu.:1.562e+08   3rd Qu.: 77844746
## Max.    :3.553e+09   Max.    :387367903
## NA's    :3668       NA's    :4482
```

Missingness


- What does this mean for "generalizability"
 - "Generalizability": Do our conclusions from this data extend ("generalize") to the population at large?

```
p <- mv %>%  
  mutate(missing = ifelse(is.na(gross) | is.na(budget),1,0)) %>%  
  group_by(year) %>%  
  summarise(propMissing = mean(missing)) %>% # Calculate the  
proportion of observations missing either gross or budget  
  ggplot(aes(x = year,y = propMissing)) +  
  geom_bar(stat = 'identity') +  
  labs(x = 'Year',y = '% Missing') +  
  scale_y_continuous(labels = scales::percent) # Format the y-axis  
Labels
```


Missingness

- We can only speak to post-2000s Hollywood!

p



Follow the process: Look

- What **type** of variables are earnings (**gross**) and costs (**budget**)?


```
mv %>%  
  drop_na(gross,budget) %>%  
  select(gross,budget) %>% glimpse()
```

```
## Rows: 3,179  
## Columns: 2  
## $ gross   <dbl> 73677478, 53278578, 723586629, 11490339, 62...  
## $ budget  <dbl> 93289619, 10883789, 160147179, 6996721, 139...
```

- Looks like continuous measures to me!

2. Univariate Visualization

```
mv %>%  
  select(title,gross,budget) %>%  
  pivot_longer(names_to = "metric",values_to = "dollars",cols =  
c("gross","budget")) %>%  
  ggplot(aes(x = dollars,color = metric)) +  
  geom_density()
```




Log and Skew

- Univariate visualization highlights significant **skew** in both measures
 - Most movies don't cost a lot and don't make a lot, but there are a few blockbusters that pull the density way out
- Let's **wrangle** two new variables that take the log of these skewed measures
 - Logging transforms skewed measures to more "normal" measures
 - This is helpful for regression!

```
mv <- mv %>%  
  mutate(gross_log = log(gross),  
         budget_log = log(budget))
```

2. Univariate Visualization


```
mv %>%  
  select(title,gross_log,budget_log) %>%  
  pivot_longer(names_to = "metric",values_to = "log_dollars",cols =  
c("gross_log","budget_log")) %>%  
  ggplot(aes(x = log_dollars,color = metric)) +  
  geom_density()
```



3. Conditional Analysis

- Continuous X continuous variables? Scatter with `geom_point()`!

```
mv %>%  
  ggplot(aes(x = budget_log, y = gross_log)) +  
  geom_point()
```



3. Conditional Analysis


- (BTW, I know I've been violating the tenets of data viz for several slides now. Let's fix that.)

```
pSimple <- mv %>%
  drop_na(budget,gross) %>%
  mutate(profitable = ifelse(gross >
budget, 'Profitable', 'Unprofitable')) %>%
  ggplot(aes(x = budget,y = gross,text = paste0(title, ' (',genre,',
',year,')')))) +
  geom_point() +
  scale_x_log10(labels = scales::dollar) +
  scale_y_log10(labels = scales::dollar) +
  labs(title = "Movie Costs and Returns",
        x = "Costs (logged budget)",
        y = "Returns (logged gross)")

pFancy <- pSimple + geom_point(aes(color = profitable)) +
  scale_color_manual(guide = 'none',values = rev(c('red','black')))
```

3. Conditional Analysis


pFancy



Look with `plotly`

- If curious, can use `plotly` to see outliers

```
require(plotly)
ggplotly(pFancy, tooltip = 'text')
```



4. Regression!

```
pSimple +  
  geom_smooth(aes(group = 1), method = 'lm', se = F)
```



4. Regression!

```
m <- lm(gross_log ~ budget_log, data = mv)
tidy(m)
```

```
## # A tibble: 2 × 5
##   term          estimate std.error statistic  p.value
##   <chr>         <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)    1.26     0.310     4.07 0.0000473
## 2 budget_log     0.964    0.0179    54.0 0
```

Interpretation

- Remember the equation: $Y = \alpha + \beta * X$
- Our Y is logged gross
- Our X is logged budget
- Thus we can rewrite as $gross_log = \alpha + \beta * budget_log$
- What is α ? What is β ?

$$gross_log = \underbrace{1.26}_{\alpha} + \underbrace{0.96}_{\beta} * budget_log$$

Interpreting with Logs

- Previously, we said:
 - α is the value of Y when X is zero
 - We need to convert back out of logged values using the `exp()` function
 - When `budget_log` is zero, the budget is `exp(0)` or \$1
- Thus, we say when the budget is \ \$1, the movie makes 1.26 logged dollars, or \$3.53

```
exp(1.26107)
```

```
## [1] 3.529196
```

Interpreting with Logs

- For the β coefficient, it depends on where the logged variable appears:
 1. $\log(Y) \sim X$: 1 unit change in $X \rightarrow (\exp(\beta) - 1) * 100$ percent change in Y
 2. $Y \sim \log(X)$: 1% increase in $X \rightarrow \beta/100$ unit change in Y
 3. $\log(Y) \sim \log(X)$: 1% increase in $X \rightarrow \beta$ percent change in Y
- In our example, a 1% increase in the budget corresponds to a 0.96% increase in gross
- You will either need to memorize these rules, or (like me) just look them up every time

Evaluation

- Every regression line makes mistakes
 - If they didn't, they wouldn't be good at **reducing complexity**!
- How bad do ours look?
 - How should we begin to answer this question!?
- Are there patterns to the mistakes?
 - We **overestimate** gross for movies that cost between \$1m and \$10m
 - These are the "indies"
 - We also **underestimate** gross to the "blockbusters"
- Why?

Understanding Regression Lines

- Regression lines choose α and β to minimize mistakes
 - Mistakes (aka "errors" or "residuals") are captured in the ε term
 - We can apply the **process** to these!


```
# Wrangle data to drop missingness!
mv_analysis <- mv %>% drop_na(gross_log,budget_log)
m <- lm(gross_log ~ budget_log,data = mv_analysis)
mv_analysis$predictions <- predict(m)
mv_analysis$errors <- mv_analysis$gross_log - mv_analysis$predictions

summary(mv_analysis$errors)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -8.2672 -0.6354   0.1648   0.0000   0.7899   8.5599
```


Univariate Viz of Errors

```
mv_analysis %>%  
  ggplot(aes(x = errors)) +  
  geom_histogram() +  
  labs(x = 'Errors: Gross - Predicted', y = 'Number of movies')
```




Univariate Viz of Errors

- Note that they are on average zero
 - Don't feel too proud! Mean 0 error is baked into the method
 - More concerned about **skew**...there is evidence of overestimating
- Can we do more? **Conditional Analysis**
 - Conditional on the **predictor** (the X variable)

Multivariate Viz of Errors

- Ideal is where errors are unrelated to predictor
 - This **should** appear as a rectangular cloud of points around zero


pIdeal



Multivariate Viz of Errors

- This is not the case for us!

```
mv_analysis %>%  
  ggplot(aes(x = budget, y = errors)) +  
  geom_point() + geom_hline(yintercept = 0, linetype = 'dashed') +  
  scale_x_log10(label = scales::dollar) + geom_smooth()
```



Multivariate Viz of Errors

- Evidence of a U-shape → underpredict low and high budgets, overpredict middle budgets
- Ergo, our model is **not great!**
 - Could add additional predictors X_2 , X_3 , etc.
 - Next lecture!

RMSE

- Univariate / Multivariate visualization of errors is **important**
- But we want to summarize model quality in a simpler way
- **RMSE**: summarizes model performance with a *single number*
 - Useful for comparing multiple models to each other

RMSE

- **Error** (ε): actual outcome (Y_i) - predicted outcome (\hat{Y}_i)
 - The "distance" between the data and the model
- **Squared**: ε^2
 1. Makes all values positive
 2. **Exaggerates** the presence of larger errors
- **Mean**: average these squared errors
- **Root**: take their square root (**un-exaggerate**)

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2}$$

RMSE

- **Error** (ϵ): actual outcome (Y_i) - predicted outcome (\hat{Y}_i)
 - The "distance" between the data and the model
- **Squared**: ϵ^2
 1. Makes all values positive
 2. Exaggerates the presence of larger errors
- **Mean**: average these squared errors
- **Root**: take their square root (un-exaggerate)

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n \underbrace{(Y_i - \hat{Y}_i)^2}_{\epsilon^2}}$$

RMSE

- **Error** (ε): actual outcome (Y_i) - predicted outcome (\hat{Y}_i)
 - The "distance" between the data and the model
- **Squared**: ε^2
 1. Makes all values positive
 2. Exaggerates the presence of larger errors
- **Mean**: average these squared errors
- **Root**: take their square root (un-exaggerate)

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n \underbrace{(\varepsilon)^2}_S}$$

RMSE

- **Error** (ε): actual outcome (Y_i) - predicted outcome (\hat{Y}_i)
 - The "distance" between the data and the model
- **Squared**: ε^2
 1. Makes all values positive
 2. Exaggerates the presence of larger errors
- **Mean**: average these squared errors
- **Root**: take their square root (un-exaggerate)

$$RMSE = \sqrt{\underbrace{\frac{1}{n} \sum_{i=1}^n (SE)}_M}$$

RMSE

- **Error** (ε): actual outcome (Y_i) - predicted outcome (\hat{Y}_i)
 - The "distance" between the data and the model
- **Squared**: ε^2
 1. Makes all values positive
 2. Exaggerates the presence of larger errors
- **Mean**: average these squared errors
- **Root**: take their square root (un-exaggerate)

$$RMSE = \sqrt{(MSE)}$$

RMSE

- RMSE is a **single measure that summarizes model performance**

```
e <- mv_analysis$gross_log - mv_analysis$predictions
se <- e^2
mse <- mean(se)
rmse <- sqrt(mse)
# Or
(rmseBudget <- sqrt(mean(mv_analysis$errors^2)))
```

```
## [1] 1.280835
```

- Is this good?

Predicting with uncertainty

- Say we're talking to investors about a new movie that costs \$10m
 - How do we plug 10m into our model?

```
summary(m)$coefficients
```

```
##              Estimate Std. Error  t value Pr(>|t|)  
## (Intercept) 0.03776401 0.03153666  1.197464 0.2314101  
## X           2.07852434 0.03152731 65.927738 0.0000000
```

- $\hat{Y}_i = \alpha + \beta * X$
 - $\alpha = 1.26$ and $\beta = 0.96$
 - where \hat{Y}_i is predicted gross (log) and X is \$10m budget (log)

```
pred_gross_log <- 1.26 + 0.96*log(1e7)
```

Predicted Gross

- Again, convert back out of logged values with `exp()`

```
scales::dollar(exp(pred_gross_log))
```

```
## [1] "$18,501,675"
```

- Cool! We'll make \$8.5m!
 - But we know our model isn't perfect
 - Need to adjust for it's errors via **RMSE**

Incorporating RMSE

- Simple idea: add and subtract RMSE from this prediction

```
pred_gross_log_ub <- 1.26 + 0.96*log(1e7) + rmseBudget  
pred_gross_log_lb <- 1.26 + 0.96*log(1e7) - rmseBudget  
scales::dollar(exp(c(pred_gross_log_ub, pred_gross_log_lb)))
```

```
## [1] "$66,599,457" "$5,139,861"
```

- So we'll either make a \$56m profit or we'll lose almost \$5m?
- **CONCLUSION PART 2:** maybe our model isn't very good?

Introducing Cross Validation

- We ran a model on the full data and calculated the RMSE
- But this approach risks "overfitting"
 - **Overfitting** is when we get a model that happens to do well on our specific data, but isn't actually that useful for predicting elsewhere.
 - "Elsewhere": Other periods, other movies, other datasets
- **Theory:** Why care about **external validity**?
 - What is the point of measuring relationship if they don't generalize?

Introducing Cross Validation

- In order to avoid **overfitting**, we want to "train" our model on one part of the data, and then "test" it on a different part of the data.
 - Model "can't see" the test data → better way to evaluate performance
- Cross Validation: randomly split our data into a train set and test set
 - *Similar to bootstrapping*

Introducing Cross Validation (CV)

```
set.seed(1021)
# Sample our data WITHOUT replacement
train <- mv_analysis %>%
  sample_n(size = round(nrow(.)*.5),
           replace = F)

# New function...remove all the rows that are the same as train
test <- mv_analysis %>%
  anti_join(train)
```

```
## Joining with `by = join_by(title, rating, genre, year,
## released, score, votes, director, writer, star, country,
## budget, gross, company, runtime, id, imdb_id,
## bechdel_score, boxoffice_a, language, gross_log,
## budget_log, predictions, errors)`
```

- We now have two datasets of roughly the same number of observations, but none of them are the same!

CV to Calculate RMSE

- We want to estimate a model based on the **test** data
- And evaluate RMSE based on the **train** data

```
m2 <- lm(gross_log ~ budget_log,train)

# predict() function on a new dataset
test$preds <- predict(m2,newdata = test)

# Now calculate RMSE on the new dataset
e <- test$gross_log - test$preds
se <- e^2
mse <- mean(se,na.rm=T)
rmse <- sqrt(mse)
rmse
```

```
## [1] 1.28959
```

CV to Calculate RMSE

- We did worse with CV! This is a *feature*
 - We are not being overconfident
 - We are avoiding "overfitting"
- Want to do this many times (like bootstrapping)

CV to Calculate RMSE

```
set.seed(123)
bsRes <- NULL
for(i in 1:100) {
  # Create training dataset
  train <- mv_analysis %>%
    sample_n(size = round(nrow(.)*.5),
              replace = F)

  # Create test dataset
  test <- mv_analysis %>%
    anti_join(train)

  mTrain <- lm(gross_log ~ budget_log, train)


  test$preds <- predict(mTrain, newdata = test)

  rmse <- sqrt(mean((test$gross_log - test$preds)^2, na.rm=T))
  bsRes <- c(bsRes, rmse)
}
mean(bsRes)
```

```
## [1] 1.279899
```

CV to Calculate RMSE

```
data.frame(rmseBS = bsRes) %>%  
  ggplot(aes(x = rmseBS)) +  
  geom_density() +  
  geom_vline(xintercept = mean(bsRes), linetype = 'dashed')
```



Cross Validation

- In this example, we used a 50-50 split
- Often, data scientists prefer an 80-20 split
 - **Improves** the model (80% of the data is more to learn from)...
 - ...but still **protects** against overfitting

```
train <- mv_analysis %>%  
  sample_n(size = round(nrow(mv_analysis)*.8),  
            replace = F)
```

